

# QAA Assignment

Jeremy Jacobson

9/6/2021

## Talapas Modules and Setup

Modules installed in conda environment ‘QAA’ in the following order using “conda install \_\_\_\_\_”

```
fastqc/0.11.5
easybuild
STAR 2.7.9a
numpy 1.13.1
pysam 0.13.0-intel-2017b-Python-3.6.3
matplotlib 2.0.1-Python-3.6.1
HTSeq 0.9.1-Python-3.6.1 - installed from pip
```

*All work on talapas was done in an interactive node using:*

```
srun -account=bgmp -partition=bgmp -nodes=1 -ntasks-per-node=1 -time=1-0:00:00 -cpus-per-task=1 -pty bash
```

*All scripts/commands were timed using “**usr/bin/time -v**”*

*All scripts were batched in talapas using:*

```
sbatch -account=bgmp -partition=bgmp -time=1-0:00:00 -cpus-per-task=1 ./_____
```

*In the case of trimmomatic and STAR, 8 CPUs were used for multithreading functionality.*

## Read Quality Score Distributions

15\_3C\_mbnl\_S11\_L008 Read 1




Figure 1: Per Base Quality Score Distribution





Figure 2: Per Base Quality Score



## 15\_3C\_mbnl\_S11\_L008 Read 2




Figure 3: Per Base Quality Score Distribution

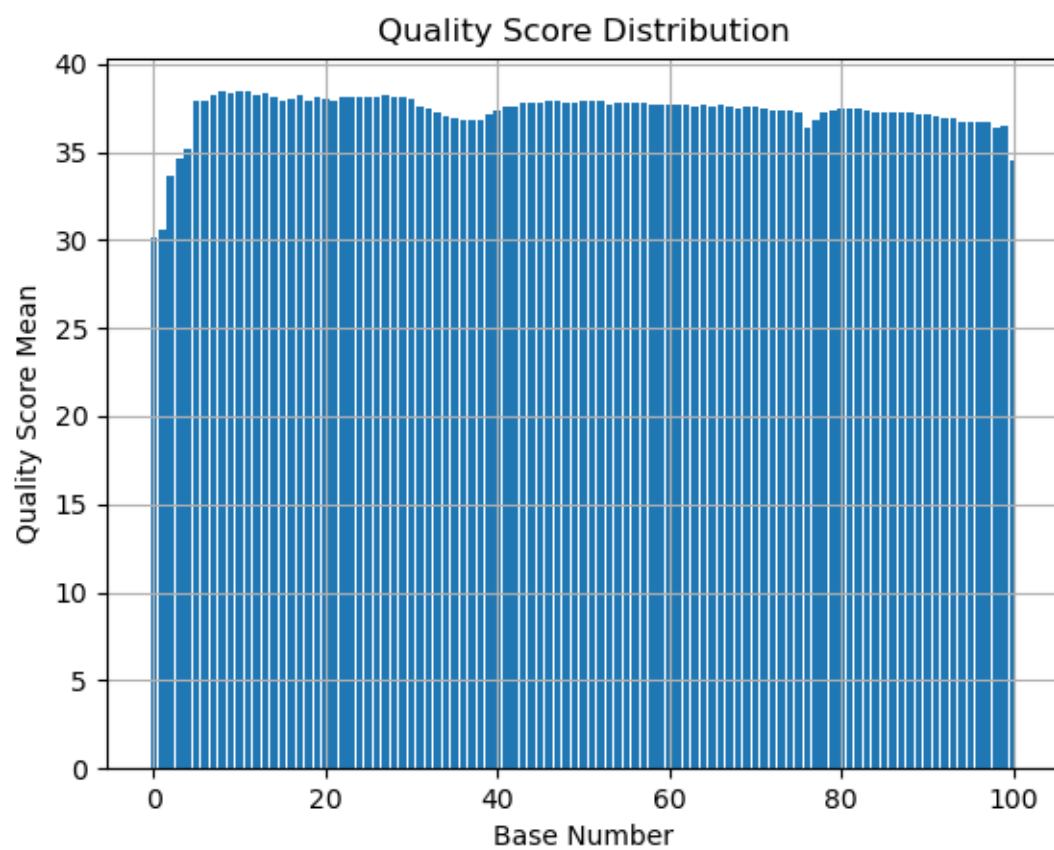



Figure 4: **Per Base Quality Score**



## 24\_4A\_control\_S18\_L008 Read 1




Figure 5: Per Base Quality Score Distribution







Figure 6: Per Base Quality Score



24\_4A\_control\_S18\_L008 Read 2




Figure 7: Per Base Quality Score Distribution




Figure 8: **Per Base Quality Score**




Figure 9: **Per Base N Content**

### Question 1

All paired graphs are consistent with each other. Increased N content is correlated with a decreased quality score. This is most visually distinct in base position 1.

Example code:

```
fastqc /projects/bgmp/shared/2017_sequencing/demultiplexed/15_3C_mbnl_S11_L008_R1_001.fastq
-o .
```

**Python ‘Per Base N Content’ graphs were generated using Demultiplexing\_Quality\_score\_plotter.py**  
 “-R\_file” argument was used to choose file. Output file was manually changed.

Python Code:

```
#!/usr/bin/env python
import Bioinfo
import matplotlib
import numpy as np
import matplotlib.pyplot as plt
import argparse

#argparse
parser = argparse.ArgumentParser()
parser.add_argument("-R_file", default=1)

args = parser.parse_args()
file = args.R_file
```

```

#Reads with 101 characters
phred_list = []
fred_list = Bioinfo.populate_list(file)
phred_list = fred_list[0]
#sums_list = phred_list[:]
line_count = fred_list[1]

print("phred_list output:", phred_list)
print("line_count output:", line_count)

count = 0
for value in phred_list:
    phred_list[count] = value/(line_count/4)
    count += 1

plt.bar(range(101), phred_list)
plt.title('Quality Score Distribution')
plt.xlabel('Base Number')
plt.ylabel('Quality Score Mean')
plt.grid(True)
plt.savefig("/home/jjacobso/bgmp/bioinformatics/Bi623/Assignments/QAA/24_R2.png")
plt.show()

```

## Question 2

The graphs generated using part of the demultiplexing script display nearly the same information as those produced through fastqc. The main difference is that fastqc provides a range of scores and a trendline. The python graphs align almost perfectly to the trendline.

Runtime:

- Fastqc: 1minute 30seconds.
- Python: 2minutes 45seconds.

Fastqc is written in java which is a compiled language and generally faster than python. It was able to generate numerous figures in the time it took python to produce one.

## Question3

Both libraries follow the same pattern. The quality scores near the earliest base positions are the lowest, hovering at 32 for R1 and 30 for R2. The rest of the qscores are around 39/40 for R1 and 37/38 for R2. The last nucleotide of each library and read also has lower quality. R1 is higher quality than R2 for both libraries because it was subjected to less chemicals and had less time to degrade before sequencing.

## Adapter Trimming Comparison

### Pt. 1: Cutadapt

Example code:

```

cutadapt -b AGATCGGAAGAGCACACGTCTGAACTCCAGTCA -o ./15_R1.fastq
/projects/bgmp/shared/2017_sequencing/demultiplexed/15_3C_mbnl_S11_L008_R1_001.fastq
Average Runtime: 1minute 30seconds Results:
15_3C_mbnl_S11_L008_R1_001:
Adapter:AGATCGGAAGAGCACACGTCTGAACTCCAGTCA -Trimmed 554754x (7.1%)
15_3C_mbnl_S11_L008_R2_001:

```

Adapter:AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT -Trimmed 582968x (7.5%)  
 24\_4A\_control\_S18\_L008\_R1\_001:  
 Adapter:AGATCGGAAGAGCACACGTCTGAACTCCAGTCA -Trimmed 489040x (4.7%)  
 24\_4A\_control\_S18\_L008\_R2\_001:  
 Adapter:AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT -Trimmed 598531x (5.7%)

## Pt. 2: Trimmomatic

Example code:

```
trimmomatic PE -threads 8 /home/jjacobso/bgmp/bioinformatics/Bi623/Assignments/QAA/15_R1.fastq
/home/jjacobso/bgmp/bioinformatics/Bi623/Assignments/QAA/15_R2.fastq
15_R1_001_trimmomatic.fastq.gz 15_R1_001_untrimmomatic.fastq.gz 15_R2_001_trimmomatic.fastq.gz
15_R2_001_untrimmomatic.fastq.gz LEADING:3 TRAILING:3 SLIDINGWINDOW:5:15 MINLEN:35
Average Runtime: 4minutes 37seconds
```

### Trimmed Reads Results Library: 15\_3C\_mbnl\_S11\_L008\_R1/R2\_001\_fastq

Input Read Pairs: 7806403 Both Surviving: 7418603 (95.03%) Forward Only Surviving: 377796 (4.84%)

Reverse Only Surviving: 5705 (0.07%) Dropped: 4299 (0.06%)

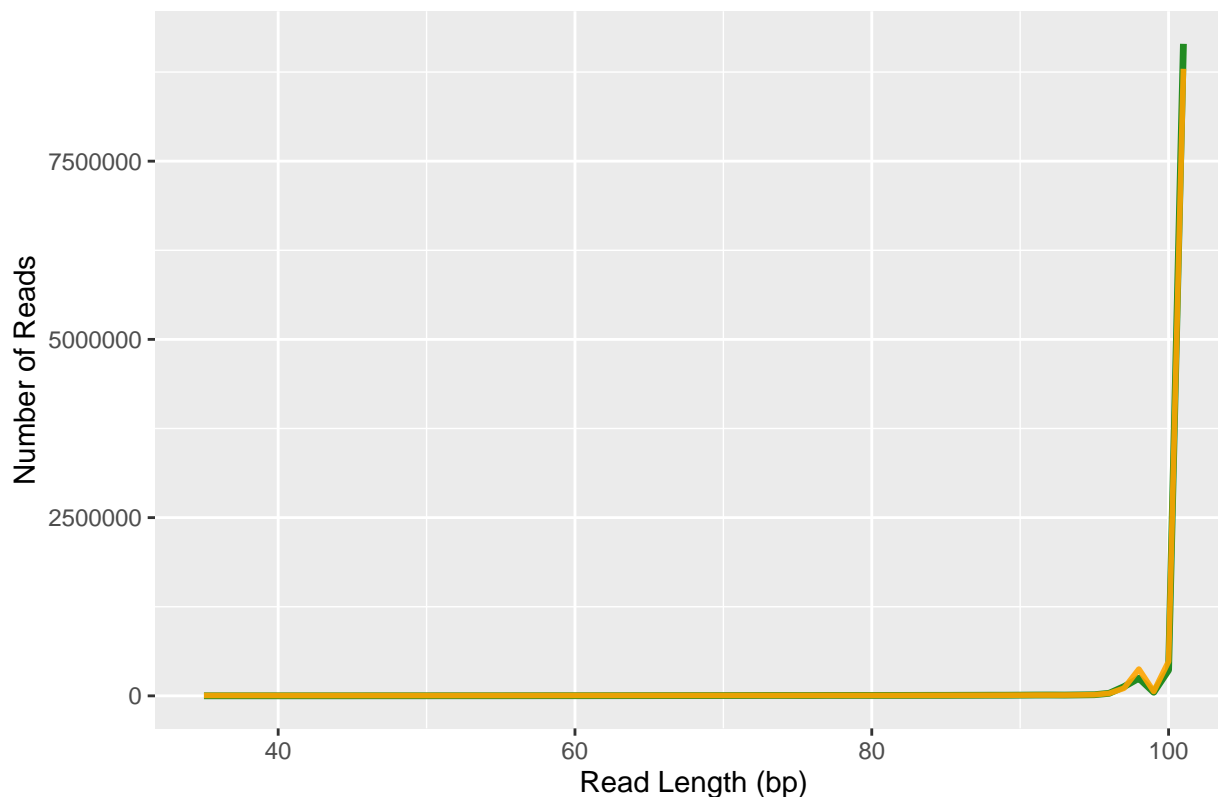
### Library: 24\_4A\_control\_S18\_L008\_R1/R2\_001\_fastq

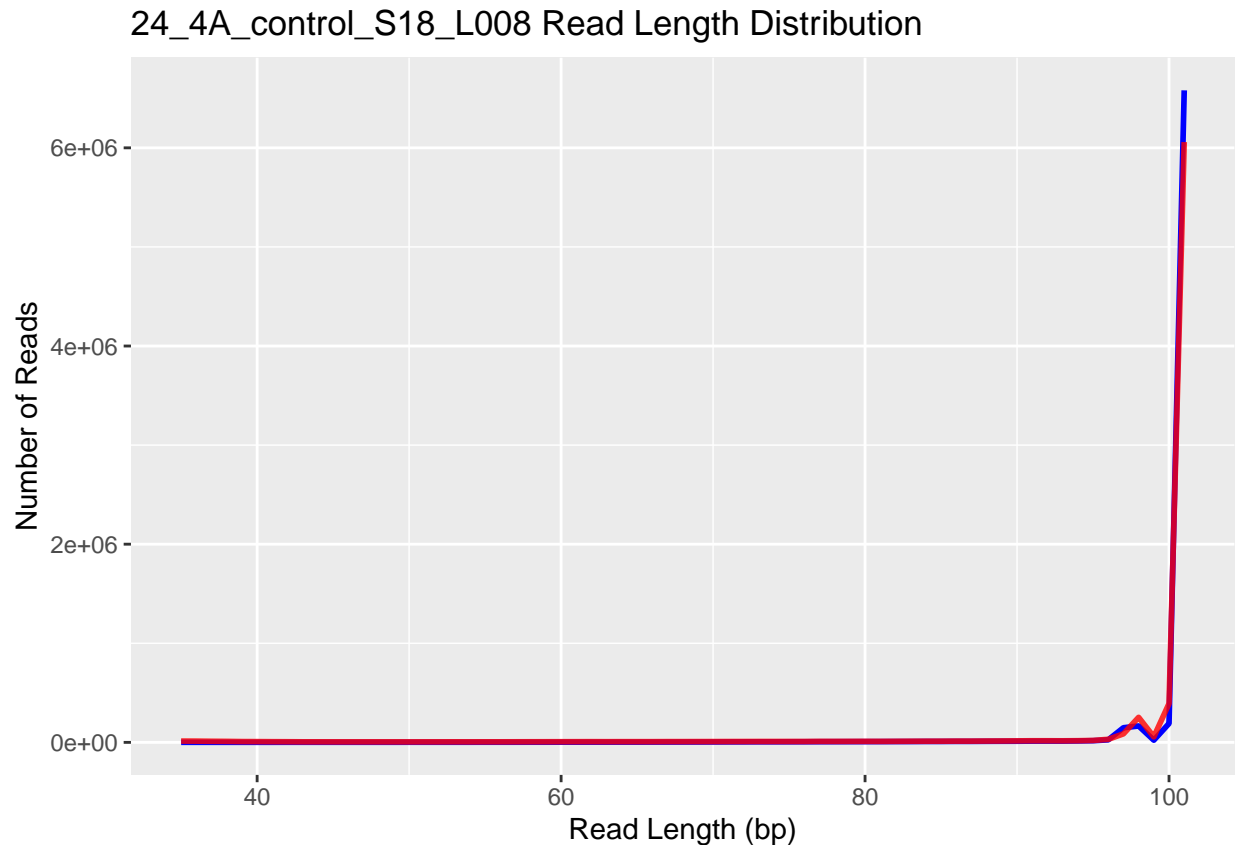
Input Read Pairs: 10515874 Both Surviving: 10245586 (97.43%) Forward Only Surviving: 255904 (2.43%)

Reverse Only Surviving: 10860 (0.10%) Dropped: 3524 (0.03%)

### Trimmed Read Distribution Plot

15\_3C\_mbnl\_S11\_L008 Read Length Distribution





**Legend:** Green: 15\_3C\_mbnl\_S11\_L008\_R1 | Red: 15\_3C\_mbnl\_S11\_L008\_R2  
 Blue: 24\_4A\_control\_S18\_L008\_R1 | Orange: 24\_4A\_control\_S18\_L008\_R2

In each paired library, the second read (red and yellow) had a decreased proportion of full length (101bp) sequences. Read 2 was likely trimmed more frequently due to lower quality scores resulting from increased exposure time / chemical degradation. Read 2 likely had more adapter trimming as well. These plots also demonstrate that library 24\_4A\_control\_S18\_L008 has fewer overall reads than 15\_3C\_mbnl\_S11\_L008.

## Alignment and Strand-Specificity

### STAR Assembly:

Example Code:

```
STAR --runThreadN 8 --runMode genomeGenerate --genomeDir /projects/bgmp/jjacobso/bioinformatics
--genomeFastaFiles /projects/bgmp/jjacobso/bioinformatics/Bi623/Assignments/QAA/Alignment_stuff/
--sjdbGTFfile /projects/bgmp/jjacobso/bioinformatics/Bi623/Assignments/QAA/Alignment_stuff/Mu
Runtime: 19minutes 17 seconds
```

### STAR Alignment:



Example code:

```
STAR --runThreadN 8 --runMode alignReads --outFilterMultimapNmax 3 --
outSAMunmapped Within KeepPairs --alignIntronMax 1000000 --alignMatesGapMax
1000000 --readFilesCommand zcat --readFilesIn /home/jjacobso/bgmp/bioinformatics/Bi623/Assignmen
/home/jjacobso/bgmp/bioinformatics/Bi623/Assignments/QAA/15_R2_001_trimmo.fastq.gz
--genomeDir /projects/bgmp/jjacobso/bioinformatics/Bi623/Assignments/QAA/Alignment_stuff/Mus
--outFileNamePrefix Mus_musculus_15_R2_001
```

Runtime: 1minute 16seconds

#### **Mus\_musculus\_15\_001\_Aligned.out.sam**

Mapped Reads: 14432097 (97.3%)

Unmapped Reads: 405109 (2.7%)

#### **Mus\_musculus\_24\_001\_Aligned.out.sam**

Mapped Reads: 19778684 (96.6%)

Unmapped Reads: 712488 (3.4%)

*Mapped and Unmapped reads were found using Map\_reader.py*

```
#!/usr/bin/python3.6
```

```
alignment_tracker = {}
```

```
mapped = 0
```

```
unmapped = 0
```

```
with open ("/projects/bgmp/jjacobso/bioinformatics/Bi623/Assignments/QAA/Alignment_stuff/Mus_musculus_24_001A
"r") as fh:
```

```
for line in fh:
```

```
if not line.startswith("@"):
```

```
line = line.split("\t")
```

```
flag = (int(line[1]))
```

```
if((flag & 4) != 4) and ((flag & 256) != 256):
```

```
mapped +=1
```

```
else:
```

```
if ((flag & 256) != 256):
```

```
unmapped +=1
```

```
print ("Mapped read count:", mapped)
```

```
print ("Unmapped read count:", unmapped)
```

#### **HTSeq Results**

##### **HTSeq:**

Example code:

```
htseq-count stranded=yes Mus_musculus_24_001Aligned.out.sam Mus_musculus.GRCm39.104.gtf
> Mus_musculus_24_001_stranded_Aligned.out.sam
```

##### **Proportion of mapped reads:**

Example code:

```
grep -v "\t0$" Mus_musculus_15_001_stranded_Aligned.out.sam | awk '{if ($1 ~
"ENS.*") sum+=$2; else sum_two+=$2} END {print (sum/(sum_two+sum))}'
```

## HtSeq Results

### **Mus\_musculus\_15\_001\_stranded\_Aligned.out.sam**

\_\_no\_feature 6612743  
\_\_ambiguous 6184  
\_\_too\_low\_aQual 13778  
\_\_not\_aligned 195207  
\_\_alignment\_not\_unique 325109

Proportion of mapped reads = **0.0357995**

### **Mus\_musculus\_15\_001\_unstranded\_Aligned.out.sam**

\_\_no\_feature 559675  
\_\_ambiguous 341555  
\_\_too\_low\_aQual 13778  
\_\_not\_aligned 195207  
\_\_alignment\_not\_unique 325109

Proportion of mapped reads = **0.806524**

### **Mus\_musculus\_24\_001\_stranded\_Aligned.out.sam**

\_\_no\_feature 9052611  
\_\_ambiguous 7141  
\_\_too\_low\_aQual 11069  
\_\_not\_aligned 350360  
\_\_alignment\_not\_unique 481260

Proportion of mapped reads = **0.033492**

### **Mus\_musculus\_24\_001\_unstranded\_Aligned.out.sam**

\_\_no\_feature 829369  
\_\_ambiguous 446219  
\_\_too\_low\_aQual 11069  
\_\_not\_aligned 350360  
\_\_alignment\_not\_unique 481260

Proportion of mapped reads = **0.79325**

**Final Conclusion:** I propose that these data are strand specific due to the above results. In libraries, 24\_4A\_control\_S18\_L008 and 15\_3C\_mbnl\_S11\_L008, the unstranded reads that mapped to genes were 80.6% and 79.3% respectively, likewise, only 3.6% and 3.3% of the stranded reads mapped to genes. Unstranded reads should theoretically map to 50% while stranded reads should map to either 100% or 0%. In this case, the stranded reads would likely be mapped at a high percentage to the reverse strand.